

Decoding The Relationship Between Genes.

Project Overview:

In this project, I work with a set of seven mutated gene sequences and try to reconstruct their genealogy using two different algorithmic strategies: a local approach and a global approach. The main idea is to first figure out how similar the gene sequences are to each other, and then use those similarities to infer grandparent-parent-child relationships between them.

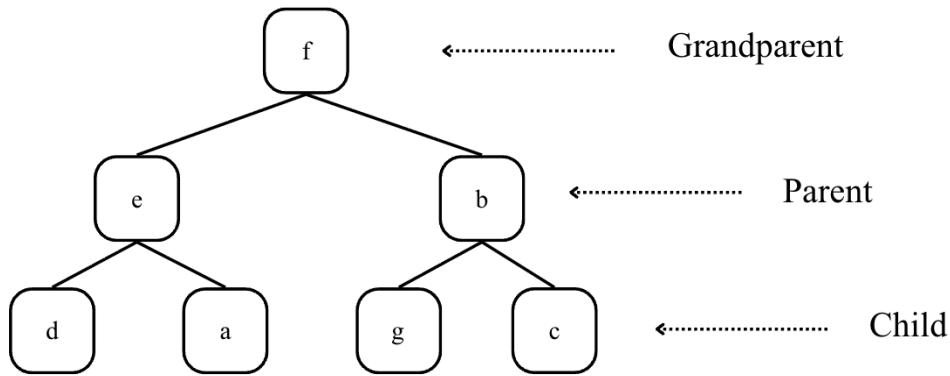


Figure 1: Inferred genealogy tree (using the global strategy)

To measure similarity, I compute the longest common subsequence (LCS) for every pair of genes and store these values in a 2D matrix. This gives me a concrete way to reason about how closely related any two sequences are.

	a	b	c	d	e	f	g
a	334	226	223	255	270	226	228
b	226	305	266	228	235	248	279
c	223	266	320	223	222	228	251
d	255	228	223	336	278	233	231
e	270	235	222	278	304	250	229
f	226	248	228	233	250	276	235
g	228	279	251	231	229	235	338

Figure 2: LCS similarity matrix (higher values indicate greater similarity).

I then convert these similarity values into distances, which I use as a proxy for how likely two genes are to be directly connected in the genealogy. Using these distances, I build two binary genealogy trees. The first uses a greedy, local strategy that makes decisions based only on immediate neighbors. The second uses a global, dynamic programming strategy that evaluates the quality of the entire tree at once. Comparing these two trees helps highlight the differences between local and global reasoning when inferring relationships.

```

local greedy tree (based on lcs distance): (5, 4, 1, 3, 0, 6, 2)
local greedy score: 0.6061
tree structure (local):
f
├── e
│   ├── d
│   └── a
└── b
    ├── g
    └── c
-----
global best tree (needleman-wunsch distance): (5, 1, 4, 2, 6, 0, 3)
global best score: -1227.0
tree structure (global):
f
├── b
│   ├── c
│   └── g
└── e
    ├── a
    └── d

```

Figure 3: Inferred genealogy trees through implementing a local and global strategy.

Finally, I go beyond tree construction by analyzing algorithmic scaling and optimization, and by using the inferred genealogies to estimate mutation probabilities, including insertions, deletions, and substitutions. This ties the algorithmic side of the project back to biological intuition and evolutionary interpretation.

```

local totals: {'ins': 185, 'del': 18, 'sub': 161, 'match': 1591}
local probs: {'p_ins': 0.09494640122511486, 'p_del': 0.009698825931597753, 'p_sub': 0.08269525267993874, 'p_match': 0.8126595201633486}

global totals: {'ins': 185, 'del': 18, 'sub': 161, 'match': 1591}
global probs: {'p_ins': 0.09494640122511486, 'p_del': 0.009698825931597753, 'p_sub': 0.08269525267993874, 'p_match': 0.8126595201633486}

```

Figure 4: Number and estimated probabilities of insertions, deletions, and substitutions.

Methods Summary:

VIDEO 1- Local and Global Strategy Explanation

Local Strategy:

I construct a binary genealogy tree using a greedy approach based on pairwise sequence similarity. I begin by converting the LCS lengths into a normalized distance matrix, where each entry is defined as $1 - \frac{LCS(i,j)}{\min(|i|, |j|)}$. The heuristic used is that smaller values indicate more similar gene sequences.

The tree is constructed through a sequence of local decisions. First, I select the grandparent node as the gene with the smallest average distance to all others, treating centrality in similarity space as a greedy heuristic for ancestral placement. Next, I choose the two parent nodes as the genes closest to the grandparent. Finally, the remaining genes are assigned as children by selecting the closest available nodes for each parent independently.

To quantify the quality of the inferred local tree, I compute a local greedy score defined as the sum of pairwise distances along all parent-child edges in the tree. Lower scores indicate tighter clustering of related genes and therefore a more consistent genealogy under the chosen distance metric.

```
local greedy tree (based on lcs distance): (5, 4, 1, 3, 0, 6, 2)
local greedy score: 0.6061
tree structure (local):
  f
  |
  +-- e
  |  |
  |  +-- d
  |  |
  |  +-- a
  |
  +-- b
      |
      +-- g
      |
      +-- c
```

Figure 5: Inferred genealogy trees through implementing a local strategy.

This strategy is local because each decision is made using only immediate pairwise distances and is never revisited. The algorithm does not evaluate alternative tree configurations or optimize a global objective. It builds the tree incrementally by choosing the best local option at each step.

Global Strategy:

I construct a binary genealogy tree by explicitly optimizing a global objective that evaluates the quality of the entire tree at once. Instead of relying only on local similarity, this approach uses Needleman-Wunsch global alignment scores to capture the entire sequence similarity between every pair of genes. These scores are converted into a pairwise distance matrix, where smaller values indicate more similar sequences.

Rather than making incremental greedy choices, I evaluate all valid binary tree configurations consistent with the fixed grandparent-parent-child structure and compute a total tree score for each candidate. The score is defined as the sum of distances along all parent-child edges in the tree. The globally optimal tree is the one that minimizes this total edge score, thereby balancing similarity across the entire genealogy rather than favoring any single local decision.

This strategy is global because the placement of each node is determined by how it affects the overall tree score. Decisions are not made independently or at every step. Instead, every possible configuration is compared under the same global metric, and the tree with the best score is selected. This ensures that the final structure reflects the strongest overall consistency among all the gene relationships.

```
global best tree (needleman-wunsch distance): (5, 1, 4, 2, 6, 0, 3)
global best score: -1227.0
tree structure (global):
  f
    └── b
      ├── c
      └── g
    └── e
      ├── a
      └── d
```

Figure 6: Inferred genealogy trees through implementing a global strategy.

Local vs Global Strategy Comparison:

As expected, the genealogy trees produced by the local and global strategies differ slightly. The local strategy builds the tree using a greedy, bottom-up process that relies only on immediate pairwise distances. Once a node is placed, the decision is never revisited, which makes the approach efficient but sensitive to early choices and unable to correct suboptimal placements.

In contrast, the global strategy evaluates the genealogy as a whole. Rather than committing to local decisions, it compares all valid tree configurations under a fixed structure and selects the one that minimizes a global objective based on Needleman-Wunsch alignment distances. This allows the algorithm to capture deeper shared structure between genes that may not be apparent from local similarity alone.

Although both strategies produce tree scores, these values cannot be compared directly, as they are calculated from different distance metrics and use different numerical scales. Instead, the comparison shows a “conceptual” trade-off: the local strategy prioritizes proximity and efficiency, while the global strategy prioritizes overall consistency across all the possible choices to create the genealogy tree.

Complexity Analysis:

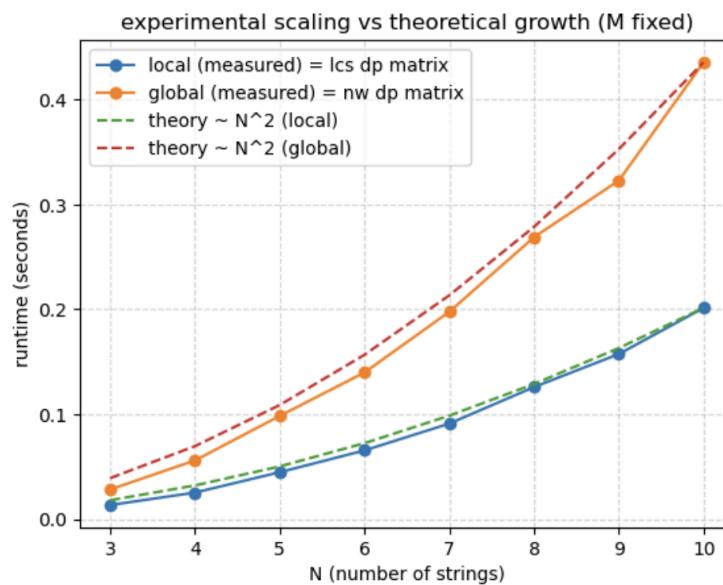


Figure 7: Experimental scaling vs. theoretical growth (M fixed).

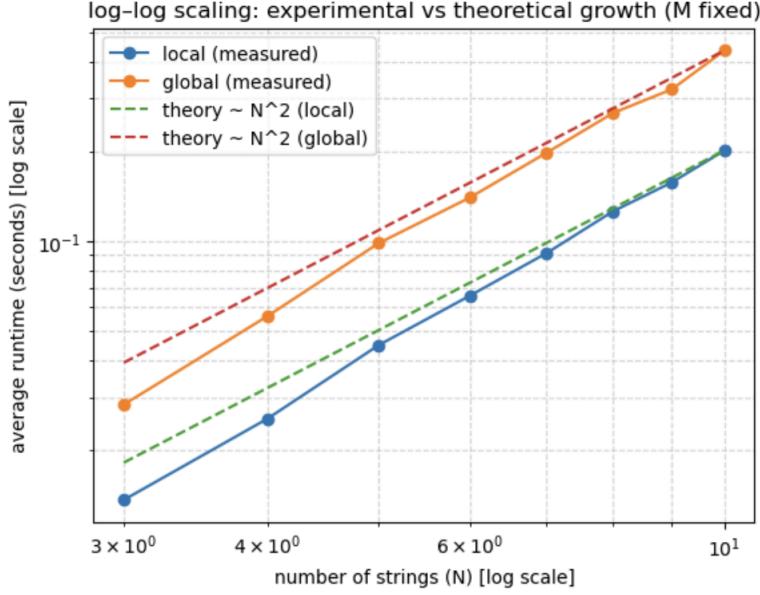


Figure 8: Log-log scaling of experimental vs. theoretical runtime (M fixed).

To study scaling, I fixed the gene length at $M = 200$ and varied the number of genes $N \in \{3, \dots, 10\}$. For each N , I generated a random dataset with `make_random_set_strings(N, M, seed = N)`, extracted the sequences with `unpack_set_strings`, and then measured average runtime using `avg_runtime_seconds(..., repeats=3)`.

For the local strategy, I timed only `lcs_dist_matrix_from_lcs_length(seqs)` inside `run_greedy()`. This isolates the dominant cost: computing pairwise LCS dynamic programming across all pairs of genes. Each LCS computation is $O(M^2)$, and there are $O(N^2)$ pairs, so the dominant runtime is $\Theta(N^2 M^2)$.

For the global strategy, I timed only `nw_dist_matrix_from_score(seqs)` inside `run_global()`. Needleman-Wunsch global alignment is also a DP over prefixes with cost $O(M^2)$ per pair, again across $O(N^2)$ pairs, so the dominant runtime is also $\Theta(N^2 M^2)$. I intentionally did not time the later tree-search step, because my Q3 tree inference was built around a fixed 7-node structure; isolating the distance-matrix DP step keeps the scaling experiment valid when N changes.

Since M is fixed, both theoretical curves reduce to $\Theta(N^2)$ growth, which is why I plotted $theory_greedy = N^2$ and $theory_global = N^2$ (scaled only for visual comparison). In both the regular and log-log plots, the measured runtimes follow the expected N^2 trend: as the number of genes increases, the runtime grows roughly quadratically. This makes sense because the dominant cost in both strategies is computing dynamic programming alignments for all pairs of sequences. The global (Needleman-Wunsch) curve consistently sits above the local (LCS) curve because each alignment is more expensive to compute, even though both methods scale the same way asymptotically.

Mutation Probability Estimation:

[Video 2](#) - Explanation on how to estimate the probabilities of insertions, deletions, and mutations given the genealogy.

Definitions:

Mutations are defined using aligned parent-child sequences. An insertion occurs when the child has a base and the parent has a gap. A deletion occurs when the parent has a base and the child has a gap. A substitution occurs when both have bases but they differ. A match occurs when both have the same base. Each aligned column corresponds to exactly one of these events.

Algorithmic Approach:

To estimate mutation probabilities, I treat each parent-child edge in the inferred genealogy as one evolutionary step. For every edge, I compute a global Needleman-Wunsch alignment and use the traceback to identify insertions, deletions, substitutions, and matches along that edge. These counts are counted across all edges in the tree to compensate for the small dataset size. Because the dataset is limited, I apply Laplace smoothing to avoid zero probabilities. Alignment ambiguity is handled using deterministic tie-breaking in the traceback, ensuring consistent and reproducible estimates.

Results and Interpretation:

The estimated probabilities show that matches dominate, while insertions and substitutions occur at moderate rates and deletions are pretty rare in comparison. This makes sense from a biological perspective: most bases are similar because genes are related, while substitutions and insertions introduce “manageable” variation. Deletions are especially “bad” because removing a base can cause a frameshift, disrupting all downstream codons and often leading to nonfunctional proteins. My hypothesis is that, from an evolutionary perspective, deletions are more strongly selected against because it creates a deleterious mutation, which can explain their lower estimated probability.

Learning for Growth-Supporting Materials:

Video 3 - LO and HC application.

Moment from class: last breakout from last session

LO application: #AlgoStratDataStruct

HC: #networks

AI Statement:

I used AI primarily as a support tool while working on this assignment. It helped me think through the overall structure of the project and debug parts of my code when I got stuck on implementation challenges. The core ideas, algorithmic choices, and biological reasoning were developed by me, but AI helped me move past specific points of confusion more efficiently by suggesting ways to organize functions, fix errors, or clarify logic. I treated AI as a way to speed up iteration and deepen my understanding, rather than as a replacement for my own problem-solving or reasoning.

References:

2.5: *The Needleman-Wunsch Algorithm*. (2020, October 5). Biology LibreTexts.

[https://bio.libretexts.org/Bookshelves/Computational_Biology/Book%3A_Computational_Biology_-_Genomes_Networks_and_Evolution_\(Kellis_et_al.\)/02%3A_Sequence_Alignment_and_Dynamic_Programming/2.05%3A_The_Needleman-Wunsch_Algorithm](https://bio.libretexts.org/Bookshelves/Computational_Biology/Book%3A_Computational_Biology_-_Genomes_Networks_and_Evolution_(Kellis_et_al.)/02%3A_Sequence_Alignment_and_Dynamic_Programming/2.05%3A_The_Needleman-Wunsch_Algorithm)

Cormen, T. H., Leiserson, C. E., Rivest, R. L., & Stein, C. (2009). *Introduction to Algorithms*, third edition. MIT Press.

CS110 Session 25 - [14.2] Synthesis of algorithms and computational applications. (2017).

Forum. Minerva.edu.

<https://forum.minerva.edu/app/courses/3717/sections/12784/classes/98193>

Interactive demo for Needleman–Wunsch algorithm. (n.d.). Bioboot.github.io.

https://bioboot.github.io/bimm143_W20/class-material/nw/

Moreanu, A. (2025a). *VIDEO 1.* Loom.

<https://www.loom.com/share/52436fdb527441778123f5f3994ac6f1>

Moreanu, A. (2025b). *Video 2.* Loom.

<https://www.loom.com/share/9771c7fa796d45159f990d32a9425a25>

Moreanu, A. (2025c). *Video 3.* Loom.

<https://www.loom.com/share/299e5c0ddac54dda9d8120d4aae0e707>

Needleman-Wunsch Explained Clearly. (2022, February 7). YouTube.

<https://www.youtube.com/watch?v=FIxYGV7WPA8>

Wikipedia Contributors. (2021, March 24). *Needleman–Wunsch algorithm*. Wikipedia;

Wikimedia Foundation.

https://en.wikipedia.org/wiki/Needleman%E2%80%93Wunsch_algorithm

Code Appendix:

Outline:

1. Problem setup + data (set_strings)
2. Q1 LCS function (all LCSs) + tests
3. Q2 LCS-length matrix + interpretation

- ```

4. Q3 Local vs Global strategy (design)
5. Q4 Implement + draw trees + compare
6. Q5 Complexity (theory + experimental plots)
7. Q6 Mutation probability estimation (idea + code + sanity check)

```

Q1 LCS function (all LCSS) + tests

```

def lcs_tables(x, y):
 """
 Construct the DP tables for the LCS problem, following the algorithm LCS-LENGTH
 from Cormen et al.

 Function compares two strings x and y and computes:
 1) A table of LCS lengths for all prefixes of the strings.
 2) A table of directional choices that records how each LCS length was obtained.

 Parameters

 x : str
 The first input string.
 y : str
 The second input string.

 Returns

 c : list of lists of int
 The LCS length table, where c[i][j] is the length of a longest common
 subsequence of the prefixes x[:i] and y[:j].
 b : list of lists of str
 The direction table used for reconstruction. Each entry is one of:
 - "NW" : diagonal move (characters matched),
 - "N" : move up (skip a character from x),
 - "W" : move left (skip a character from y).

 """
 # lenghts of strings
 m = len(x)
 n = len(y)

 # create c and b tables
 c = [[0 for _ in range(n + 1)] for _ in range(m + 1)]
 b = [['' for _ in range(n + 1)] for _ in range(m + 1)]

```

```

fill in tables
for i in range(1, m + 1):
 for j in range(1, n + 1):

 if x[i - 1] == y[j - 1]:
 c[i][j] = c[i - 1][j - 1] + 1
 b[i][j] = "NW" # diagonal

 else:
 if c[i - 1][j] >= c[i][j - 1]:
 c[i][j] = c[i - 1][j]
 b[i][j] = "N" # up
 else:
 c[i][j] = c[i][j - 1]
 b[i][j] = "W" # left

return c, b

def reconstruct_lcs(x, y, c, i, j, memoization):
 """
 Recursively reconstruct all LCSs of x[:i] and y[:j] using the CLRS LCS length
 table.

 This function follows the same logic used in Cormen et al. to reconstruct an LCS,
 but
 extends it to return all the possible LCSs.

 Memoization ensures that each subproblem (i, j) is solved only once, avoiding
 redundant
 recursion when multiple optimal paths exist

 Parameters

 x : str
 input of first string
 y : str
 input of second string
 c : list of lists of int
 LCS length table produced by teh lcs_tables()
 i : int
 the current index into x
 """
 if i == 0 or j == 0:
 return [""]
 if b[i][j] == "N":
 return [y[0:j] + s for s in reconstruct_lcs(x, y[1:j], c, i, j - 1, memoization)]
 if b[i][j] == "W":
 return [x[0:i] + s for s in reconstruct_lcs(x[1:i], y, c, i - 1, j, memoization)]
 if b[i][j] == "NW":
 return reconstruct_lcs(x, y, c, i - 1, j - 1, memoization) + [x[i-1] + y[j-1]] + reconstruct_lcs(x[0:i-1], y[0:j-1], c, i - 1, j - 1, memoization)

```

```

j : int
 the current index into y
memoization : dict
 Memoization dictionary mapping (i, j) to a set of LCS strings

Returns

set of str: all longest common subsequences of x[:i] and y[:j].
"""

if we have already solved this subproblem, we return the stored result
if (i, j) in memoization:
 return memoization[(i, j)]

base case: one of the prefixes is empty
if i == 0 or j == 0:
 memoization[(i, j)] = {""}
 return memoization[(i, j)]

case 1: last characters match
if x[i - 1] == y[j - 1]:
 previous_lcs = reconstruct_lcs(x, y, c, i - 1, j - 1, memoization)

 result = set()
 for seq in previous_lcs:
 result.add(seq + x[i - 1])

 memoization[(i, j)] = result
 return result

case 2: last characters do not match
results = set()

move up if it preserves the optimal LCS length
if c[i - 1][j] >= c[i][j - 1]:
 upper_results = reconstruct_lcs(x, y, c, i - 1, j, memoization)
 results.update(upper_results)

move left if it preserves the optimal LCS length
if c[i][j - 1] >= c[i - 1][j]:
 left_results = reconstruct_lcs(x, y, c, i, j - 1, memoization)
 results.update(left_results)

```

```

memoization[(i, j)] = results
return results

def find_lcs(x, y):
 """
 This function finds all LCSSs between two strings.
 It follows the dynamic programming approach described in Cormen et al. by first
 computing the LCS length table and then backtracking through this table to recover
 the subsequences.
 Cormen et al. describes how to recover a single LCS, but this function extends the
 method to return all possible LCSSs of maximum length.

 Parameters

 x : str
 First input string
 y : str
 Second input string

 Returns

 (list[str] | None, int)
 A tuple containing:
 - a sorted list of all longest common subsequences,
 - the length of the LCS.

 If there is no common subsequence, returns (None, 0).

 """
 # handle edge cases where one string is empty
 if not x or not y:
 return None, 0

 # building the LCS tables using the CLRS algorithm
 c, _ = lcs_tables(x, y) #we only need the length table c, the direction table b is
 not required here

 # the LCS length is stored in the bottom-right cell of the table
 lcs_length = c[len(x)][len(y)]

 if lcs_length == 0:

```

```

 return None, 0

 # reconstruct all LCSSs using the memoization
 memoization = {}
 lcs_set = reconstruct_lcs(x, y, c, len(x), len(y), memoization)

 # remove empty string if present
 lcs_set.discard("")

 return sorted(lcs_set), lcs_length

```

```

Test 1: CLRS example
x1 = "ABCBDAB"
y1 = "BDCABA"
assert find_lcs(x1, y1) == (['BCAB', 'BCBA', 'BDAB'], 4)

Test 2: one empty string
tests the base case
x2 = "abc"
y2 = ""
assert find_lcs(x2, y2) == (None, 0)

Test 3: one single character match
edge case with small input --> see if DP table is reconstructed correctly
x3 = "abc"
y3 = "a"
assert find_lcs(x3, y3) == (['a'], 1)

Test 4: case sensitivity check
edge case to check if character nature matters
x4 = "ABC"
y4 = "abc"
assert find_lcs(x4, y4) == (None, 0)

Test 5: identical strings
best case scenario for LCS --> the full string is returned; no duplicates introduced
x5 = "genetics"
y5 = "genetics"
assert find_lcs(x5, y5) == (['genetics'], len("genetics"))

Test 6: no common characters
edge case where no characters match

```

```

x6 = "xyz"
y6 = "abc"
assert find_lcs(x6, y6) == (None, 0)

Test 7: Mixed-case characters (partial match)
checking if the algorithm handles symbol matching correctly
x7 = "AbC"
y7 = "abc"
assert find_lcs(x7, y7) == (['b'], 1)

print("All LCS tests passed!")

```

Q2 LCS-length matrix + interpretation

```

set_strings = [
 ('a',
 "TGGTGGAACATTGATCTGCTTTACGGTACAGTCTAGTTTGGTACCCCTGCAGAACGATTGGGCCATCTACAGTGCCGCG
 CCACAGTTAAAGTAGTGACGTGGATCTGATATTAACAGAGGACGTTGGACGGAAACTATCACCCAGCTAACAAATCATATGA
 TGAACCCGACGTTACGGTATGAGTGACCCATTACTGTCAACACGCATGAAGATCGGCCAGCTCTAGTAATGACC
 GTTGCTCCACATATGTTCCGACATCTGTGCGTCTTGGAACCTATATCGTAGCAGGAGGGATGTACTCGG") ,

 ('b',
 "TGGTACGAAGGCATCTCTTCCCGTGGCGTATGGTGTATGGTATCCACCGAACGGTGGCGTCTACAACCTCTCCAAGATA
 TACGAGCGAGTTAGAGCTAATTAAACAGAGGGCGTCGGTACGTCATTAGGCACCCAGCATGAATCAACCGCATGAACCGTTATC
 GTGGGTTTCAGGACTCTATGTAAGGATGAAGATTTCGCAACAGCTACTAAATAATGGAGTGTACAATAATGCAACCCCTACA
 CCGTCAATCTGCACTGCTAGTAGAGCCTGGGATTGGTCCCTCGG") ,

 ('c',
 "TGGTCGAAGGCATCTTTGCCGTGGCGTGTGGTGTATTGGCAATCCACCGAACGGTGGCGTCTCGTCTACCAACTCTCCAAGAG
 TATACGAGCGCAGTCAGACCTCTACATAAACATAGGGCGTCCTTACGTCATTAGGCACCCAGCGATGAATCAACCTCATCGAAC
 CCTATGGTGGTTACGCAGCTCTACTAAAAGCACTCGTAGGTTTGTCAACAGCTACTAAATAATTCAAGGTGTGCAAATT
 GCATACCCAACACCGTGCAATCTGCTGTGATAGTAGAGCATGGGGAGTCTAGGGACCTCGG") ,

 ('d',
 "GGGTGCGACCAAGATCTATTCTCCGTACGGATAAGGTCTCTATTGTCCTGGAAATGATTGGAAACCTACATTGAGCG

```

```

AGCAGTATAAGTAGGGACGTAGATCATTATATTATCTAGAGGGCGTGGTTGGGACGCAAGTTCCAGCGGAACTAGGAATCGT
CCGATGAACTCCTACGATGAGAGAGTCCTGAACTATAGGGCAGAAGGATTGATATGCCGAGGACCACGATCATTAGTAATGA
GCGTTGCGTCACATATGTATCCCGACCATTGGTCTTGGAACTATCTTGATCGTACCCACGGAGGATGTTCTCGG") ,

('e',

"GGGTGCGACAAGATCTCTTCCGTGACGTATAGTTTATGGTACTCCCCGAAACGATTGGCAACTACAATCCCGAACAGTAT
AAGTAGTGACGTAGATCTTATATTAAACAGAGGGCGTTGGACGCAAGTATCACCAACTAGAATCATCCGATGAACCTCGACGTGA
GGTGCAGTGAACATAGGGCAAAGGATGAATATGCGCGAGGCCACGCTCTAGTAATGAGCGTTGCTCACATATGTATCCGACAT
CGGTCGTCTTGGAACTACCTAGATCGTAGCAGGAGGATGTTACTCGG") ,

('f',

"GGGTGCGAAAGCATCTCTTCCGTGGCGTATAGTTTATGGTATCCCCGAAACGCTGGCTACTACAATCTCGAAGTATAGAGT
GAGTAGATTAAACAGAGGGCGTCGTTGACGCATTAGCACCAACTGAATCAACCGATAACTAACGTGGGTTCACTGACTAT
AGGGCAAAGGATGAACATTTGAGCAGCTAAATAATGAGCGTGACAATATGAATCCACACCGTCATCTGAACCTAGATTGA
GCAAGAGGTTGTTCTCGG") ,

('g',

"GGGACGTAGGCACCTCTTCCGTGGCGTATGGGTGTGACGGGTATCCACCCGTACGGGCACCTCTAACAACGCTCTCACT
AAGATCATAACGAGCGAGTATAGAGCTTAATTACCAAGAGGGACCTCGCTGCCGACATTAGAGGCAACCGATGAACCTCAAGCCCCA
ATGTAACGTCATCGTGGGATTGCAGGACTCTATATAAAAGGATGAAAGGAATTTCCGCAACCAGTTACTAAAGTAAATGGAGTG
TGTACAAATAATGTGAAGCCCTCACACCGTGCCTCTTGTACTGCGTAGTAAGGAGCAGTTACTAAAGTAAATGGAGTG")
]

```

```

def lcs_length(x, y):
 """
 Computing the length of the LCS between two strings.

 This function follows the same dynamic programming recurrence as the
 LCS-LENGTH algorithm described in Cormen et al., but only computes
 the final LCS length (not the subsequence itself).

 In order to reduce memory usage, we store only two rows of the DP table at a time:
 the previous row and the current row.

 This is possible because each DP c[i][j] depends only on c[i-1][j], c[i][j-1], and
 c[i-1][j-1].
 Parameters

```

```

x : str
y : str

Returns

int
 Length of the longest common subsequence between x and y.

"""

making sure that y is the shorter string to minimize memory usage
if len(y) > len(x):
 x, y = y, x

m, n = len(x), len(y)

prev[j] represents c[i-1][j]
curr[j] represents c[i][j]
prev = [0] * (n + 1)
curr = [0] * (n + 1)

Build the DP table row by row (CLRS recurrence)
for i in range(1, m + 1):
 xi = x[i - 1]
 curr[0] = 0 # base case: LCS with empty prefix is 0

 for j in range(1, n + 1):
 if xi == y[j - 1]:
 # If characters match, then extend the LCS
 curr[j] = prev[j - 1] + 1
 else:
 # If characters do not match, then take the best of skipping one
 character
 curr[j] = max(prev[j], curr[j - 1])

 # move current row to prev for next iteration
 prev, curr = curr, prev

Final LCS length is in the last cell of the last completed row
return prev[n]

```

```
import numpy as np
```

```

"""
Constructing a pairwise LCS length matrix for all strings in set_strings.

Each entry len_lcs_matrix[i, j] stores the length of the LCS between the i-th and j-th
strings.

This matrix summarizes pairwise similarity between all gene sequences.

"""

n = len(set_strings)

initialize a 7x7 matrix of integers
len_lcs_matrix = np.zeros((n, n), dtype=int)

for i in range(n):
 seq_i = set_strings[i][1]

 len_lcs_matrix[i, i] = len(seq_i) # the LCS of a string with itself is the full
length of the string

 for j in range(i + 1, n):
 seq_j = set_strings[j][1] #compute only the upper triangle and mirror (matrix
is symmetric)

 # Compute LCS length using DP (CLRS)
 length = lcs_length(seq_i, seq_j)

 len_lcs_matrix[i, j] = length
 len_lcs_matrix[j, i] = length # symmetry

len_lcs_matrix

```

```

"""
Non-trivial pairs correspond to comparisons between different strings (i != j).
We are extracting only the upper-triangular entries (i < j) to avoid duplicates.

"""

upper_triangle_values = len_lcs_matrix[np.triu_indices(n, k=1)]

```

```

unique_lcs_lengths = np.unique(upper_triangle_values)

print("Unique LCS lengths (non-trivial pairs):", unique_lcs_lengths)
print("Number of unique LCS lengths:", len(unique_lcs_lengths))

```

```

"""
Printing the matrix in a readable format to manually examine it.

"""

labels = [label for label, _ in set_strings]

Print column headers
print(" " + " ".join(f"{lab:>4}" for lab in labels))

Print each row with its corresponding label
for i in range(n):
 row_values = " ".join(f"{len_lcs_matrix[i, j]:>4}" for j in range(n))
 print(f"{labels[i]:>4} {row_values}")

```

Q3 Local vs Global strategy

```

def unpack_set_strings(set_strings):
 """
 Return labels and sequences as two parallel lists.

 I use this as a small helper so the rest of my code is easier to read.
 It also helps me avoid repeating tuple indexing everywhere.

 Parameters

 set_strings : list of tuples
 A list where each element is of the form (label, sequence).

 Returns

 labels : list of str
 The list of labels (e.g. ['a', 'b', 'c', ...]).
 sequences : list of str
 The corresponding list of DNA sequences.

```

```

"""
labels = []
sequences = []

for label, seq in set_strings:
 labels.append(label)
 sequences.append(seq)

return labels, sequences

def tree_from_indices(labels, root, left_child, right_child):
 """
 Build a tiny nested-dict tree from explicit parent/child indices.

 This is a convenience function so I can print trees in a consistent way.
 """
 return {
 "name": labels[root],
 "left": {"name": labels[left_child], "left": None, "right": None},
 "right": {"name": labels[right_child], "left": None, "right": None},
 }

def print_tree(tree, indent=0):
 """
 Print a nested-dict binary tree in a readable way.
 """
 if tree is None:
 return

 print(" " * indent + tree["name"])
 print_tree(tree["left"], indent + 1)
 print_tree(tree["right"], indent + 1)

```

Local Strategy:

Metric: LCS distance from Q2 matrix

```

def lcs_distance_matrix_from_lengths(len_lcs_matrix, sequences):
 """
 Convert the LCS length matrix into a normalized distance matrix in [0, 1].

 Each distance is computed as:
 1 - (LCS length / min(sequence lengths))

 This is suitable for a local strategy because it is a simple pairwise metric based
 on the information we already computed in Q2.

 Heuristic = smaller distance means more similar

 Parameters

 len_lcs_matrix : np.ndarray
 A square matrix where entry (i, j) is the LCS length between sequences i and j.
 sequences : list of str
 The DNA sequences corresponding to the matrix indices.

 Returns

 np.ndarray
 A square distance matrix with values in the range [0, 1].
 """
 n = len(sequences)
 dist = np.zeros((n, n), dtype=float)

 for i in range(n):
 for j in range(n):
 denom = min(len(sequences[i]), len(sequences[j]))
 dist[i, j] = 1.0 - (len_lcs_matrix[i, j] / denom)

 return dist

```

Greedy tree building (closest neighbors)

```

def argmin_average_distance(dist):
 """
 Find the index with the smallest average distance to all other nodes.

```

```

This function is used as a greedy heuristic to select a root-like node
(grandparent) that is centrally located with respect to all others.

Parameters

dist : np.ndarray
 A square distance matrix.

Returns

int
 The index of the node with the smallest average distance.

"""
n = dist.shape[0]
best_i = 0
best_avg = float("inf")

for i in range(n):
 avg = (np.sum(dist[i]) - dist[i, i]) / (n - 1)
 if avg < best_avg:
 best_avg = avg
 best_i = i

return best_i

def two_closest_nodes(dist, anchor, candidates):
 """
 Identify the two candidate nodes closest to a given anchor node (smallest distance)
 => these will be parents (most "similar" to the grandparent).

 This function supports greedy local decisions by comparing only immediate neighbors
 based on pairwise distance.

 Parameters

 dist : np.ndarray
 A square distance matrix.
 anchor : int
 Index of the reference node.
 candidates : list of int
 Indices of candidate nodes to compare against the anchor.
 """

```

```

 Returns

 tuple of int
 The indices of the two closest candidate nodes.

 """
 # store (distance, node) pairs and pick the smallest two
 pairs = []
 for node in candidates:
 pairs.append((dist[anchor], node))

 pairs.sort(key=lambda t: t[0])
 return pairs[0][1], pairs[1][1]

def local_greedy_tree_from_distance(dist, labels):
 """
 Construct a genealogy tree (grandparent -> 2 parents -> 4 children) using a local
 greedy strategy.

 Logic:
 1) Select a central node with minimal average distance as the grandparent.
 2) Choose the two closest nodes to act as parents.
 3) Assign remaining nodes to parents based on closest local distance.

 This strategy is local because each decision considers only immediate neighbor
 comparisons, not the global quality of the entire tree.

 Parameters

 dist : np.ndarray
 A square distance matrix between all sequences.
 labels : list of str
 Labels corresponding to each node index.

 Returns

 tuple of int
 A 7-tuple representing the tree structure:
 (grandparent, parent1, parent2, child11, child12, child21, child22)
 """

```

```

n = dist.shape[0]
nodes = list(range(n))

choose grandparent (most central by average distance)
g = argmin_average_distance(dist)

choose two closest nodes to g as the parents
remaining = [x for x in nodes if x != g]
p1, p2 = two_closest_nodes(dist, g, remaining)

remove parents from the pool
remaining = [x for x in remaining if x not in (p1, p2)]

assign children to each parent using a greedy local rule
i keep this simple: each parent takes its two closest remaining nodes
c11, c12 = two_closest_nodes(dist, p1, remaining)
remaining = [x for x in remaining if x not in (c11, c12)]

c21, c22 = two_closest_nodes(dist, p2, remaining)

return (g, p1, p2, c11, c12, c21, c22)

```

Tree Score (to help compare local vs global strategy)

```

def tree_edge_score(dist, tree_tuple):
 """
 Compute the total edge score for a fixed 7-node genealogy tree.

 The score is the sum of distances along all parent-child edges.
 Lower scores indicate a tighter clustering of related nodes.

 Parameters

 dist : np.ndarray
 A square distance matrix.
 tree_tuple : tuple of int
 A tuple encoding the tree structure as node indices.

 Returns
 """

```

```

float
 The total edge distance score for the tree.
"""
g, p1, p2, c11, c12, c21, c22 = tree_tuple

score = 0.0
score += dist[g, p1]
score += dist[g, p2]
score += dist[p1, c11]
score += dist[p1, c12]
score += dist[p2, c21]
score += dist[p2, c22]

return score

```

Global Strategy:

Metric: Needleman-Wunsch (DP)

```

def needleman_wunsch_score(x, y, match=1, mismatch=-1, gap=-1):
 """
 Compute the Needleman-Wunsch global alignment score between two sequences.

 This function uses dynamic programming, following the same prefix-based recurrence
 style as CLRS. Each DP cell represents the best alignment score for two prefixes of
 the sequences.

 Unlike LCS, this method enforces end-to-end alignment and therefore captures global
 similarity between sequences.

 Parameters

 x : str
 First DNA sequence.
 y : str
 Second DNA sequence.
 match : int
 Score for matching characters.

```

```

mismatch : int
 Penalty for mismatched characters.

gap : int
 Penalty for inserting a gap.

>Returns

int
 The optimal global alignment score.

"""
m, n = len(x), len(y)
dp = [[0 for _ in range(n + 1)] for _ in range(m + 1)]

base cases: aligning against empty prefix forces gaps
for i in range(1, m + 1):
 dp[i][0] = dp[i - 1][0] + gap
for j in range(1, n + 1):
 dp[0][j] = dp[0][j - 1] + gap

for i in range(1, m + 1):
 xi = x[i - 1]
 for j in range(1, n + 1):
 yj = y[j - 1]

 diag = dp[i - 1][j - 1] + (match if xi == yj else mismatch)
 up = dp[i - 1][j] + gap
 left = dp[i][j - 1] + gap

 dp[i][j] = max(diag, up, left)

return dp[m][n]

```

Build a NW distance matrix + choose the best tree globally (search)

```

def nw_distance_matrix(sequences, match=1, mismatch=-1, gap=-1):
 """
 Build a pairwise distance matrix using Needleman-Wunsch alignment scores.

```

```
Higher alignment scores indicate greater similarity, so distances are computed as
the negative of the alignment score.
```

```
Parameters

sequences : list of str
 DNA sequences to compare.
match : int
 Score for matching characters.
mismatch : int
 Penalty for mismatched characters.
gap : int
 Penalty for gaps.

Returns

np.ndarray
 A symmetric distance matrix based on global alignment.
"""
n = len(sequences)
dist = np.zeros((n, n), dtype=float)

for i in range(n):
 for j in range(i, n):
 if i == j:
 dist[i, j] = 0.0
 else:
 s = needleman_wunsch_score(sequences[i], sequences[j],
 match=match, mismatch=mismatch, gap=gap)
 dist[i, j] = -float(s)
 dist[j, i] = dist[i, j]

 return dist

def best_tree_by_global_score(dist):
 """
 Identify the best genealogy tree by minimizing a global distance objective.

 This function evaluates all valid tree configurations under a fixed
 grandparent-parent-child structure and selects the one with the lowest total edge
 distance.
 """
```

```
Because the entire tree is evaluated at once => global strategy.
```

```
Parameters
```

```

```

```
dist : np.ndarray
```

```
 A square distance matrix.
```

```
Returns
```

```

```

```
best_score : float
```

```
 The minimum total edge distance found.
```

```
best_tree : tuple of int
```

```
 The tree structure that achieves this minimum score.
```

```
"""
```

```
n = dist.shape[0]
```

```
nodes = list(range(n))
```

```
best_score = float("inf")
```

```
best_tree = None
```

```
choose grandparent
```

```
for g in nodes:
```

```
 others = [x for x in nodes if x != g]
```

```
choose two parents (unordered pair)
```

```
for i in range(len(others)):
```

```
 for j in range(i + 1, len(others)):
```

```
 p1 = others[i]
```

```
 p2 = others[j]
```

```
kids = [x for x in others if x not in (p1, p2)]
```

```
choose 2 kids for p1, remaining 2 go to p2
```

```
for a in range(len(kids)):
```

```
 for b in range(a + 1, len(kids)):
```

```
 c11 = kids[a]
```

```
 c12 = kids[b]
```

```
 remaining = [x for x in kids if x not in (c11, c12)]
```

```
 c21 = remaining[0]
```

```
 c22 = remaining[1]
```

```

 candidate = (g, p1, p2, c11, c12, c21, c22)
 score = tree_edge_score(dist, candidate)

 if score < best_score:
 best_score = score
 best_tree = candidate

 return best_score, best_tree

```

Q4 Implement + draw trees + compare

```

labels, sequences = unpack_set_strings(set_strings)

local strategy uses your lcs matrix from q2
lcs_dist = lcs_distance_matrix_from_lengths(len_lcs_matrix, sequences)
local_tree_tuple = local_greedy_tree_from_distance(lcs_dist, labels)
local_score = tree_edge_score(lcs_dist, local_tree_tuple)

print("local greedy tree (based on lcs distance):", local_tree_tuple)
print("local greedy score:", round(local_score, 4))

g, p1, p2, c11, c12, c21, c22 = local_tree_tuple
print("tree structure (local):")
print(labels[g])
print("├", labels[p1])
print("| ├", labels[c11])
print("| └", labels[c12])
print("└", labels[p2])
print(" ├", labels[c21])
print(" └", labels[c22])

print("\n" + "—" * 80 + "\n")

global strategy uses needleman-wunsch distance
nw_dist = nw_distance_matrix(sequences, match=1, mismatch=-1, gap=-1)
best_score, best_tree_tuple = best_tree_by_global_score(nw_dist)

print("global best tree (needleman-wunsch distance):", best_tree_tuple)

```

```

print("global best score:", round(best_score, 4))

g, p1, p2, c11, c12, c21, c22 = best_tree_tuple
print("tree structure (global):")
print(labels[g])
print("├-", labels[p1])
print("| └-", labels[c11])
print("| └-", labels[c12])
print("└-", labels[p2])
print(" ├-", labels[c21])
print(" └-", labels[c22])

```

## Q5 Complexity (theory + experimental plots)

```

import random

def random_dna_string(m):
 """
 generate a random dna string of length m.
 """
 alphabet = ["A", "C", "G", "T"]
 return "".join(random.choice(alphabet) for _ in range(m))

def make_random_set_strings(n, m, seed=0):
 """
 create a list like set_strings but with random sequences.

 Parameters

 n : int
 number of genes (strings)
 m : int
 length of each gene
 seed : int
 seed so results are reproducible

 Returns

 """

```

```

list[tuple[str, str]]
 [('a', 'ACGT...'), ('b', 'TGCA...'), ...]
"""
random.seed(seed)
labels = [chr(ord("a") + i) for i in range(n)]
return [(labels[i], random_dna_string(m)) for i in range(n)]

```

```

def lcs_dist_matrix_from_lcs_length(sequences):
 """
 build a normalized lcs distance matrix using the lcs_length function we already
 have.

 distance(i, j) = 1 - lcs_length(i, j) / min(len(i), len(j))

 Returns

 np.ndarray
 n x n symmetric distance matrix in [0, 1]
 """
 n = len(sequences)
 dist = np.zeros((n, n), dtype=float)

 for i in range(n):
 dist[i, i] = 0.0
 for j in range(i + 1, n):
 L = lcs_length(sequences[i], sequences[j])
 denom = min(len(sequences[i]), len(sequences[j]))
 d = 1.0 - (L / denom)
 dist[i, j] = d
 dist[j, i] = d

 return dist

```

```

def nw_dist_matrix_from_score(sequences, match=1, mismatch=-1, gap=-1):
 """
 build a distance matrix using needleman-wunsch scores.

 i convert score -> distance by using dist = -score.
 this keeps ordering consistent: smaller distance means more similar.

```

```

 Returns

 np.ndarray
 n x n symmetric distance matrix
 """
 n = len(sequences)
 dist = np.zeros((n, n), dtype=float)

 for i in range(n):
 dist[i, i] = 0.0
 for j in range(i + 1, n):
 s = needleman_wunsch_score(sequences[i], sequences[j],
 match=match, mismatch=mismatch, gap=gap)
 d = -float(s)
 dist[i, j] = d
 dist[j, i] = d

 return dist

```

```

import time
import matplotlib.pyplot as plt

def avg_runtime_seconds(run_fn, repeats=3):
 """
 Time a function a few times and return the average runtime.

 Parameters

 run_fn : callable
 a function that takes no arguments and runs the computation we want to time
 repeats : int
 how many times to repeat the timing (we average to reduce noise)

 Returns

 float
 average runtime in seconds
 """
 times = []
 for _ in range(repeats):

```

```

t0 = time.perf_counter()
run_fn()
t1 = time.perf_counter()
times.append(t1 - t0)
return sum(times) / len(times)

choose M fixed and vary N
M = 200
N_values = [3, 4, 5, 6, 7, 8, 9, 10]

greedy_times = []
global_times = []

for N in N_values:
 test_set = make_random_set_strings(N, M, seed=N)
 _, seqs = unpack_set_strings(test_set)

 def run_greedy():
 """
 local strategy timing.

 for q5 i only time the dominant cost of the local method:
 building the pairwise lcs distance matrix (dp over all pairs).
 the actual greedy tree building step is lower-order and also
 assumes a fixed 7-node structure in our q3 implementation,
 so timing just the dp part is the clean scaling experiment.
 """
 _ = lcs_dist_matrix_from_lcs_length(seqs)

 def run_global():
 """
 global strategy timing.

 for q5 i only time the dominant cost of the global method:
 building the pairwise needleman-wunsch distance matrix (dp over all pairs).
 the tree search/inference step is a separate layer on top of the matrix and
 in our project code it was designed specifically for 7 nodes, so this keeps
 the scaling experiment valid when n changes.
 """
 _ = nw_dist_matrix_from_score(seqs)

```

```

greedy_times.append(avg_runtime_seconds(run_greedy, repeats=3))
global_times.append(avg_runtime_seconds(run_global, repeats=3))

---- theoretical curves (scaled to match the plot shape) ----
since M is fixed, the dp part scales like:
pairwise comparisons = O(N^2)
and each comparison costs O(M^2), which is constant here
Ns = np.array(N_values, dtype=float)

theory_greedy = Ns**2
theory_global = Ns**2 # same dominant dp scaling when M is fixed

scale curves so they sit near the measured curves (for shape comparison only)
theory_greedy = theory_greedy / theory_greedy.max() * max(greedy_times)
theory_global = theory_global / theory_global.max() * max(global_times)

---- plot ----
plt.figure()
plt.plot(N_values, greedy_times, marker="o")
plt.plot(N_values, global_times, marker="o")
plt.plot(N_values, theory_greedy, linestyle="--")
plt.plot(N_values, theory_global, linestyle="--")

plt.xlabel("N (number of strings)")
plt.ylabel("runtime (seconds)")
plt.title("experimental scaling vs theoretical growth (M fixed)")
plt.legend([
 "local (measured) = lcs dp matrix",
 "global (measured) = nw dp matrix",
 "theory ~ N^2 (local)",
 "theory ~ N^2 (global)"
])
plt.grid(True, linestyle="--", alpha=0.5)
plt.show()

```

```

---- log-log plot (both axes on a log scale) ----

plt.figure()

```

```

experimental measurements
plt.loglog(N_values, greedy_times, marker="o")
plt.loglog(N_values, global_times, marker="o")

theoretical curves
plt.loglog(N_values, theory_greedy, linestyle="--")
plt.loglog(N_values, theory_global, linestyle="--")

plt.xlabel("number of strings (N) [log scale]")
plt.ylabel("average runtime (seconds) [log scale]")
plt.title("log-log scaling: experimental vs theoretical growth (M fixed)")

plt.legend([
 "local (measured)",
 "global (measured)",
 "theory ~ N^2 (local)",
 "theory ~ N^2 (global)"
])

plt.grid(True, which="both", linestyle="--", alpha=0.5)
plt.show()

```

## Q6 Mutation probability estimation

```

def nw_align(parent, child, match=1, mismatch=-1, gap=-1):
 """
 Global alignment (needleman-wunsch) with traceback.

 I need the alignment, not just the score, because i want to count: insertions,
 deletions, and substitutions along a parent -> child edge.

 parameters

 parent : str
 the "parent" gene sequence
 child : str
 the "child" gene sequence
 match, mismatch, gap : int
 scoring settings for the dp recurrence

```

```

 returns

 (str, str)
 aligned_parent, aligned_child using '-' for gaps
 """
 m = len(parent)
 n = len(child)

 # dp table of best scores for prefixes
 dp = [[0] * (n + 1) for _ in range(m + 1)]
 # traceback table: "D" diagonal, "U" up, "L" left
 move = [[None] * (n + 1) for _ in range(m + 1)]

 # initialize borders (only gaps possible)
 for i in range(1, m + 1):
 dp[i][0] = dp[i - 1][0] + gap
 move[i][0] = "U"
 for j in range(1, n + 1):
 dp[0][j] = dp[0][j - 1] + gap
 move[0][j] = "L"

 # fill dp table (same idea as clrs lcs-length, just different recurrence)
 for i in range(1, m + 1):
 for j in range(1, n + 1):
 diag = dp[i - 1][j - 1] + (match if parent[i - 1] == child[j - 1] else
mismatch)
 up = dp[i - 1][j] + gap
 left = dp[i][j - 1] + gap

 best = max(diag, up, left)
 dp[i][j] = best

 # deterministic tie-break so results are reproducible
 if best == diag:
 move[i][j] = "D"
 elif best == up:
 move[i][j] = "U"
 else:
 move[i][j] = "L"

 # traceback

```

```

aligned_p = []
aligned_c = []
i, j = m, n

while i > 0 or j > 0:
 step = move[i][j]
 if step == "D":
 aligned_p.append(parent[i - 1])
 aligned_c.append(child[j - 1])
 i -= 1
 j -= 1
 elif step == "U":
 aligned_p.append(parent[i - 1])
 aligned_c.append("-")
 i -= 1
 else: # "L"
 aligned_p.append("-")
 aligned_c.append(child[j - 1])
 j -= 1

aligned_p.reverse()
aligned_c.reverse()
return "".join(aligned_p), "".join(aligned_c)

```

```

def count_events(aligned_parent, aligned_child):
 """
 Count insertions/deletions/substitutions/matches from an alignment.

 Definitions i use for q6:
 - insertion: parent has '-' and child has a base
 - deletion: parent has a base and child has '-'
 - substitution: both have bases but they differ
 - match: both have bases and they are the same

 parameters

 aligned_parent : str
 aligned_child : str

 returns

```

```

dict
 counts for 'ins', 'del', 'sub', 'match'
"""
counts = {"ins": 0, "del": 0, "sub": 0, "match": 0}

for a, b in zip(aligned_parent, aligned_child):
 if a == "-" and b != "-":
 counts["ins"] += 1
 elif a != "-" and b == "-":
 counts["del"] += 1
 elif a != "-" and b != "-":
 if a == b:
 counts["match"] += 1
 else:
 counts["sub"] += 1

return counts

```

```

def estimate_probs_from_tree(tree_tuple, sequences, match=1, mismatch=-1, gap=-1,
alpha=1):
"""
Estimating probabilities of insertion/deletion/mutation using the genealogy.

I treat each edge as one evolutionary step (parent -> child).
For each edge I:
1) compute a global alignment using needleman-wunsch (dp + traceback)
2) count events from the alignment
then I sum counts over all edges and convert to probabilities.

parameters

tree_tuple : tuple
 (g, p1, p2, c11, c12, c21, c22) as indices into sequences
sequences : list[str]
 sequences in the same index order as the tree
match, mismatch, gap : int
 nw scoring settings
alpha : float
 laplace smoothing (helps because dataset is tiny)

returns

```

```

(dict, dict)
 (total_counts, probabilities)
"""
g, p1, p2, c11, c12, c21, c22 = tree_tuple
edges = [(g, p1), (g, p2), (p1, c11), (p1, c12), (p2, c21), (p2, c22)]

total = {"ins": 0, "del": 0, "sub": 0, "match": 0}

for parent_i, child_i in edges:
 parent_seq = sequences[parent_i]
 child_seq = sequences[child_i]

 aligned_p, aligned_c = nw_align(parent_seq, child_seq, match=match,
mismatch=mismatch, gap=gap)
 counts = count_events(aligned_p, aligned_c)

 # add into totals
 for k in total:
 total[k] += counts[k]

convert counts to probabilities with laplace smoothing
denom = (total["ins"] + total["del"] + total["sub"] + total["match"]) + 4 * alpha

probs = {
 "p_ins": (total["ins"] + alpha) / denom,
 "p_del": (total["del"] + alpha) / denom,
 "p_sub": (total["sub"] + alpha) / denom,
 "p_match": (total["match"] + alpha) / denom,
}

return total, probs

```

```

we already have these from q4
labels, sequences = unpack_set_strings(set_strings)

local_counts, local_probs = estimate_probs_from_tree(local_tree_tuple, sequences)
global_counts, global_probs = estimate_probs_from_tree(best_tree_tuple, sequences)

print("local totals:", local_counts)
print("local probs:", local_probs)

```

```
print("\nglobal totals:", global_counts)
print("global probs:", global_probs)
```